

Motivation

Context: Why the problem matters in one or two sentences.
Problem: What is expensive, slow, or brittle about the current pipeline.
Our solution: What you built, what it supervises / predicts, and what you release.




One-line caption for the teaser figure.

Methodology

Pipeline

Two or three components, each with a bold lead-in:

- **Stage A:** sampling, labels, metadata. Include the key distribution or range.
- **Stage B:** simulation / rendering / feature extraction. Dataset size and hardware.



Model


Architectures compared, loss, optimizer, augmentation, train time.

- Fine-tuned for ...with ...loss, Adam ($\text{lr} = 10^{-4}$)
- Augmentation: ...
- Trained N epochs on ...

Experimental Results

Accuracy

One paragraph: best model, headline metric, how close the others are.



Predicted vs. ground truth on the test split.

Efficiency

Speed / parameter trade-off and why it matters for deployment.

Table 1. Model comparison. Replace numbers with your results.

Model	Score	GPU FPS	CPU FPS	Params (M)
Model A	0.00	00.0	00.0	00.0
Model B	0.00	00.0	00.0	00.0
Model C	0.00	00.0	00.0	00.0

Conclusion & Future Work

Contribution: what is open-sourced and what labels / outputs it provides.
Results: best accuracy vs. best speed/accuracy trade-off, plus venue if already published.
Future work: domain gap, real data, edge deployment.

Open-Source Resources

Paper and code. Drop `qr_paper.png` / `qr_github.png` into `images/` and replace the placeholders.



Paper



Source Code